Benchmark cost vs agreement with other benchmarks (demo, data fetched 2026-09-13) | alignment on top-20 shared models by consensus

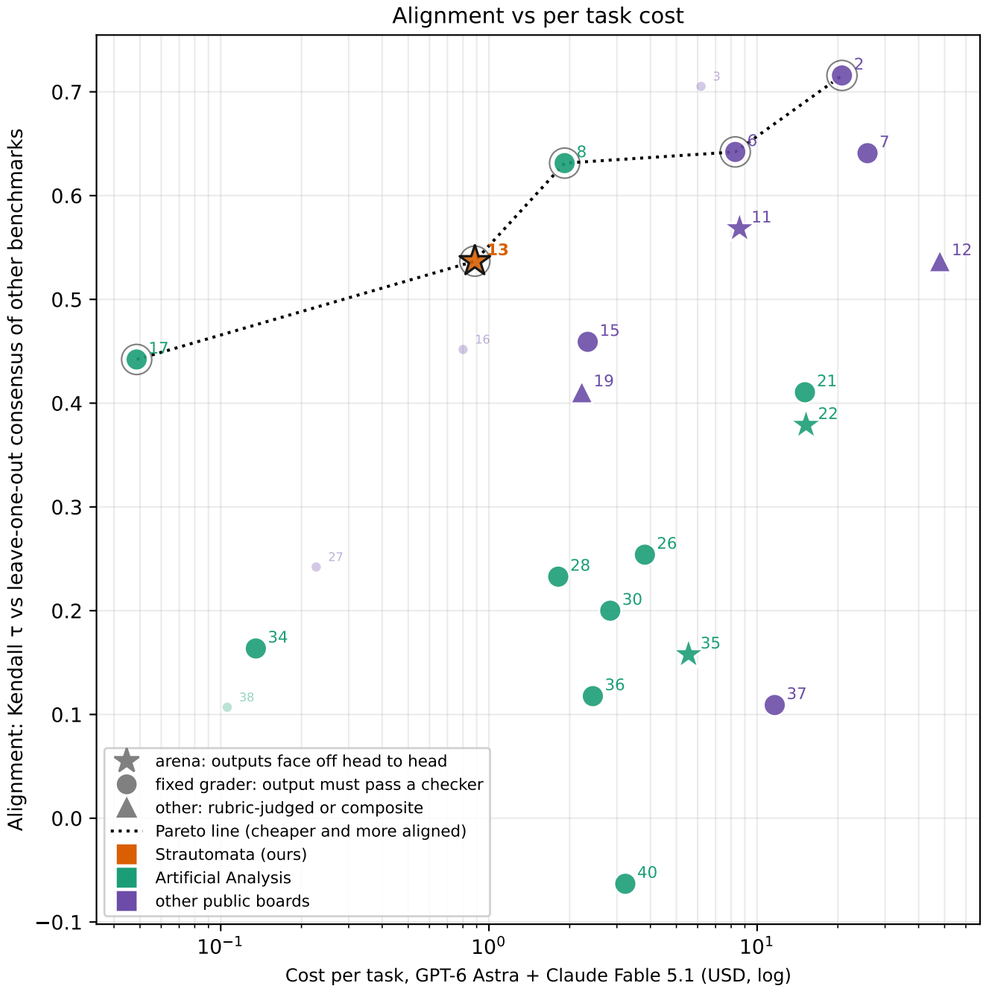

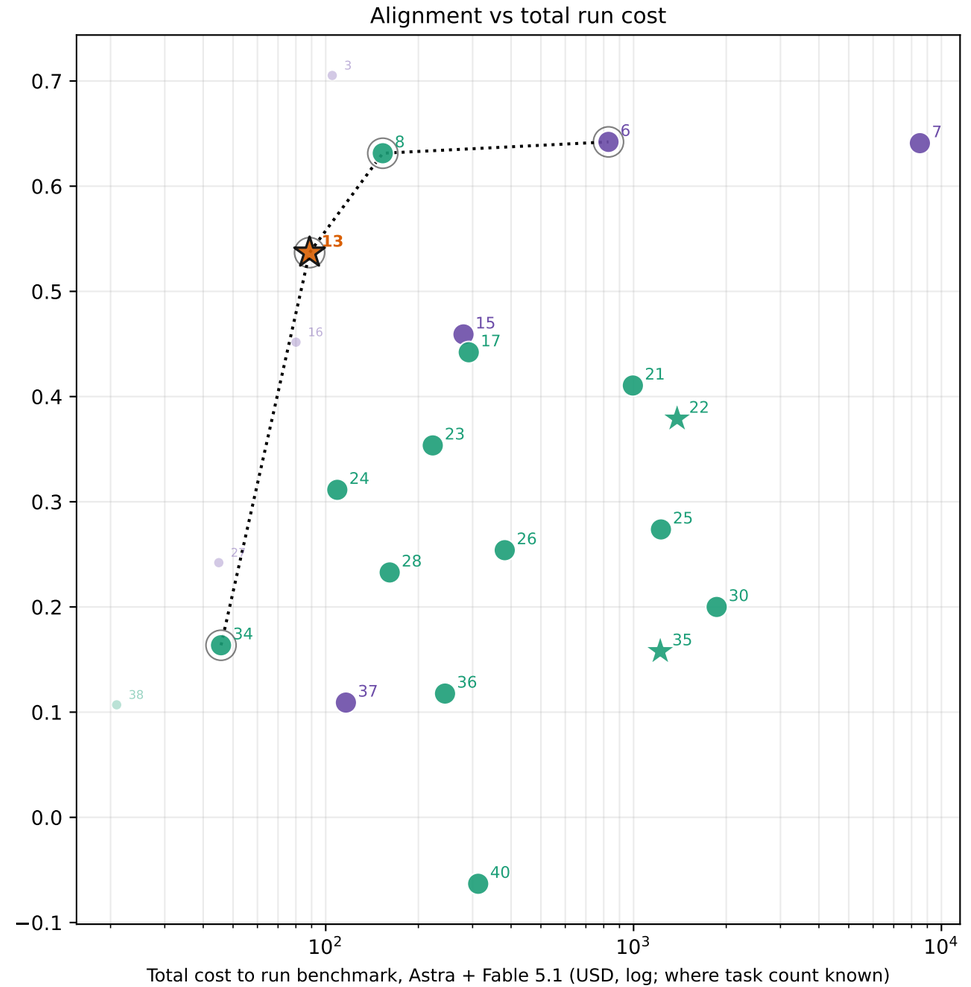

1. aa_mmlu_pro	τ=0.91 n=20	mixed			
2. matharena	τ=0.72 n=20	max	$20.72/task		
3. weirdml	τ=0.71 n=20	high	$6.18/task	SAT	
4. aa_aime25	τ=0.69 n=20	mixed			SAT
5. aa_livecodebench	τ=0.67 n=20	mixed			
6. frontiercode	τ=0.64 n=20	high	$8.29/task		
7. terminal_bench_4	τ=0.64 n=14	mixed	$25.80/task		
8. aa_analyst_agent	τ=0.63 n=20	max	$1.92/task		
9. aa_global_mmlu_lite	τ=0.60 n=14	mixed			SAT
10. aa_apex_agents	τ=0.57 n=20	mixed			
11. agent_arena	τ=0.57 n=20	max	$8.60/task		
12. vals_index	τ=0.54 n=20	max	$48.00/task		
13. Strautomata	τ=0.54 n=20	high	$0.89/task		
14. arc_agi_3	τ=0.51 n=14	mixed			
15. arc_agi_2	τ=0.46 n=20	high	$2.34/task		
16. arc_agi_1	τ=0.45 n=20	high	$0.80/task	SAT	
17. aa_omniscience	τ=0.44 n=20	high	$0.05/task		
18. aa_harvey_lab	τ=0.43 n=20	mixed			SAT
19. aicd	τ=0.41 n=20	medium	$2.22/task		
20. aa_itbench	τ=0.41 n=20	mixed			
21. aa_terminalbench_v40	τ=0.41 n=20	high	$15.07/task		
22. aa_briefcase	τ=0.38 n=20	high	$15.22/task		
23. aa_critpt	τ=0.35 n=20	high			
24. aa_mlcr	τ=0.31 n=20	max			
25. aa_hle	τ=0.27 n=20	high			
26. aa_gdp_pdf	τ=0.25 n=20	high	$3.82/task		
27. or_gpqa_diamond	τ=0.24 n=20	default	$0.23/task	SAT	
28. aa_terminalbench_v21	τ=0.23 n=20	high	$1.81/task		
29. aa_tau2	τ=0.22 n=20	mixed			SAT
30. aa_automationbench	τ=0.20 n=20	high	$2.83/task		
31. aa_terminalbench_hard	τ=0.20 n=20	mixed			
32. aa_ifbench	τ=0.19 n=20	mixed			
33. aa_enterprise_ops_gym	τ=0.18 n=20	mixed			
34. aa_scicode	τ=0.16 n=20	high	$0.14/task		
35. aa_gdpval	τ=0.16 n=20	high	$5.55/task		
36. aa_lcr	τ=0.12 n=20	high	$2.44/task		
37. real_swe	τ=0.11 n=8	default	$11.63/task		
38. aa_gpqa_diamond	τ=0.11 n=20	high	$0.11/task	SAT	
39. aa_mmmu_pro	τ=-0.01 n=20	mixed			
40. aa_tau3_banking	τ=-0.06 n=20	high	$3.22/task		

n = models shared with the consensus; effort column = reasoning effort used for the cost of Astra / Fable 5.1. consensus = mean scaled rank over the other benchmarks, min 5 shared models. Unnumbered benchmarks in the list have no cost for both models.
SAT / tiny dim marker = saturated: top-5 models all >90% and within 5 points.